

HW1

Zhaohua Chunyu

2023-02-13

In this exercise, we predict the sale price of a house using its other characteristics.

```
train = read_csv("/Users/zozochunyu/Documents/DSII/HW/DSII_HW1/housing_training.csv") %>%
  janitor::clean_names()

## Rows: 1440 Columns: 26
## -- Column specification -----
## Delimiter: ","
## chr (4): Overall_Qual, Kitchen_Qual, Fireplace_Qu, Exter_Qual
## dbl (22): Gr_Liv_Area, First_Flr_SF, Second_Flr_SF, Total_Bsmt_SF, Low_Qual_...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
test = read_csv("/Users/zozochunyu/Documents/DSII/HW/DSII_HW1/housing_training.csv") %>%
  janitor::clean_names()

## Rows: 1440 Columns: 26
## -- Column specification -----
## Delimiter: ","
## chr (4): Overall_Qual, Kitchen_Qual, Fireplace_Qu, Exter_Qual
## dbl (22): Gr_Liv_Area, First_Flr_SF, Second_Flr_SF, Total_Bsmt_SF, Low_Qual_...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
# delete rows containing the missing data
train = na.omit(train)
test = na.omit(test)

xtrain = model.matrix(sale_price ~ ., train)[,-1]
ytrain = train$sale_price

xtest = model.matrix(sale_price ~ ., test)[,-1]
ytest = test$sale_price

ctrl1 = trainControl(method = "repeatedcv", number = 10, repeats = 5)
```

Least squares

```
set.seed(2023)
lm.fit <- train(xtrain,ytrain,
               method = "lm",
               trControl = ctrl1)
```


```
pred.lm = predict(lm.fit, newx = xtest)
mse.lm = mean((ytest-pred.lm)^2)
mse.lm
```

```
## [1] 479188190
```

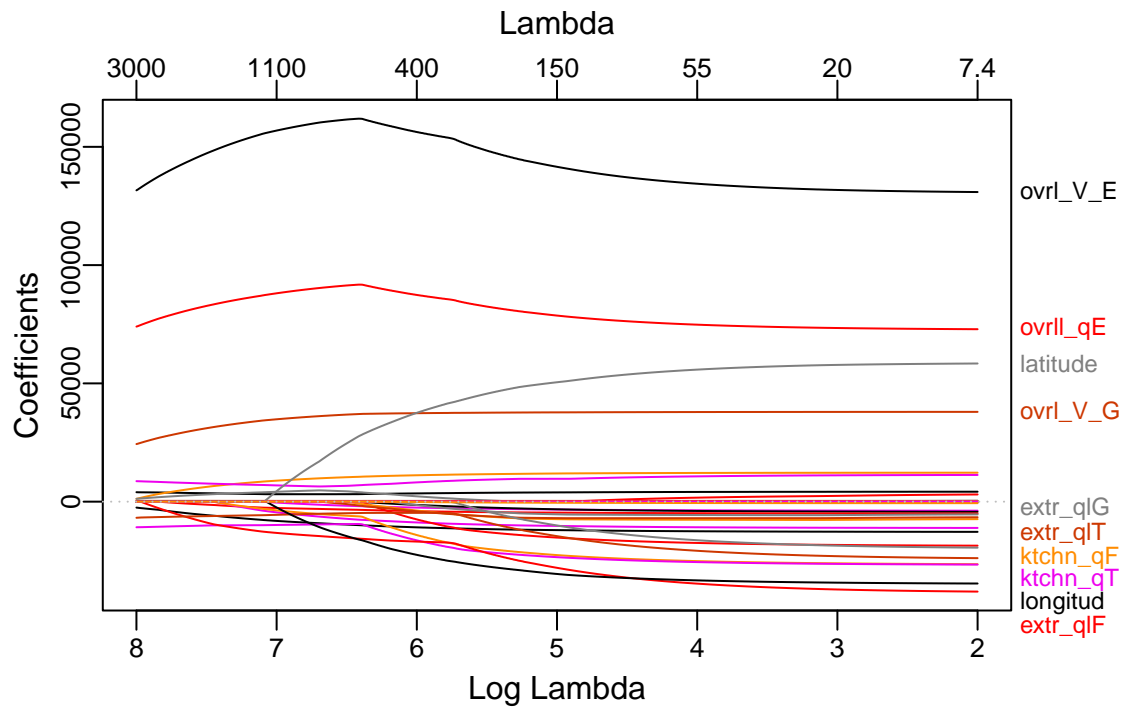
When fitting a least squares model, the test error is 4.7918819×10^8 .

LASSO

```
set.seed(2023)
cv.lasso = cv.glmnet(xtrain, ytrain,
                     standardize = TRUE,
                     alpha = 1,
                     lambda = exp(seq(8, 2, length = 1000)))
plot(cv.lasso)
abline(h = (cv.lasso$cvm + cv.lasso$cvstd)[which.min(cv.lasso$cvm)], col = 4, lwd = 2)
```



```
# cv.lasso$glmnet.fit is a fitted glmnet object using the full training data
# plot(cv.lasso$glmnet.fit, xvar = "lambda", label=TRUE)
plot_glmnet(x = cv.lasso$glmnet.fit)
```



```
cv.lasso$lambda.min
```

```
## [1] 75.06229
```

```
cv.lasso$lambda.1se
```

```
## [1] 886.0619
```

```
lasso.fit.min = predict(cv.lasso, s = "lambda.min", type = "coefficients") ; lasso.fit.min
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##               lambda.min
## (Intercept)    -4.806120e+06
## gr_liv_area      6.524898e+01
## first_flr_sf     8.136255e-01
## second_flr_sf      .
## total_bsmt_sf     3.543974e+01
## low_qual_fin_sf   -4.071470e+01
## wood_deck_sf      1.157547e+01
## open_porch_sf     1.529046e+01
## bsmt_unf_sf       -2.087834e+01
## mas_vnr_area      1.093989e+01
## garage_cars        4.064961e+03
## garage_area       8.207621e+00
## year_built        3.229258e+02
## tot_rms_abv_grd   -3.582772e+03
## full_bath         -3.779630e+03
## overall_qualAverage -4.824402e+03
## overall_qualBelow_Average -1.240562e+04
## overall_qualExcellent 7.565602e+04
## overall_qualFair   -1.069627e+04
## overall_qualGood    1.208675e+04
## overall_qualVery_Excellent 1.359728e+05
```

```
## overall_qualVery_Good      3.785440e+04
## kitchen_qualFair           -2.458207e+04
## kitchen_qualGood           -1.696340e+04
## kitchen_qualTypical        -2.509214e+04
## fireplaces                  1.043171e+04
## fireplace_quFair           -7.617540e+03
## fireplace_quGood            .
## fireplace_quNo_Fireplace    1.266012e+03
## fireplace_quPoor           -5.599848e+03
## fireplace_quTypical        -6.985489e+03
## exter_qualFair             -3.319039e+04
## exter_qualGood             -1.494540e+04
## exter_qualTypical          -1.941089e+04
## lot_frontage                9.924786e+01
## lot_area                    6.041190e-01
## longitude                   -3.270618e+04
## latitude                     5.462173e+04
## misc_val                     8.124489e-01
## year_sold                   -5.504641e+02
```

```
lasso.fit.1se = predict(cv.lasso, s = "lambda.1se", type = "coefficients") ; lasso.fit.1se
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                      lambda.1se
## (Intercept)          -1.948645e+06
## gr_liv_area           5.615314e+01
## first_flr_sf          1.142665e+00
## second_flr_sf         .
## total_bsmt_sf         3.677563e+01
## low_qual_fin_sf       -2.489128e+01
## wood_deck_sf          8.189861e+00
## open_porch_sf         7.534159e+00
## bsmt_unf_sf           -1.922749e+01
## mas_vnr_area          1.429866e+01
## garage_cars           3.145401e+03
## garage_area           1.130120e+01
## year_built            3.154739e+02
## tot_rms_abv_grd       -1.046696e+03
## full_bath             .
## overall_qualAverage    -2.969075e+03
## overall_qualBelow_Average -8.857011e+03
## overall_qualExcellent   8.971786e+04
## overall_qualFair       -5.830308e+03
## overall_qualGood        9.599970e+03
## overall_qualVery_Excellent 1.593488e+05
## overall_qualVery_Good   3.579178e+04
## kitchen_qualFair       -4.826778e+03
## kitchen_qualGood        .
## kitchen_qualTypical    -9.694576e+03
## fireplaces             6.544737e+03
## fireplace_quFair        .
## fireplace_quGood        4.607899e+03
## fireplace_quNo_Fireplace .
## fireplace_quPoor        .
## fireplace_quTypical     .
```

```
## exter_qualFair          -1.419105e+04
## exter_qualGood          .
## exter_qualTypical       -5.223227e+03
## lot_frontage            6.761224e+01
## lot_area                5.522474e-01
## longitude               -8.528497e+03
## latitude                1.364228e+04
## misc_val                .
## year_sold               .

pred.lasso.min = predict(cv.lasso, s = "lambda.min", newx = xtest)
mse.lasso.min = mean((ytest - pred.lasso.min)^2)
mse.lasso.min
```

```
## [1] 479950966
```

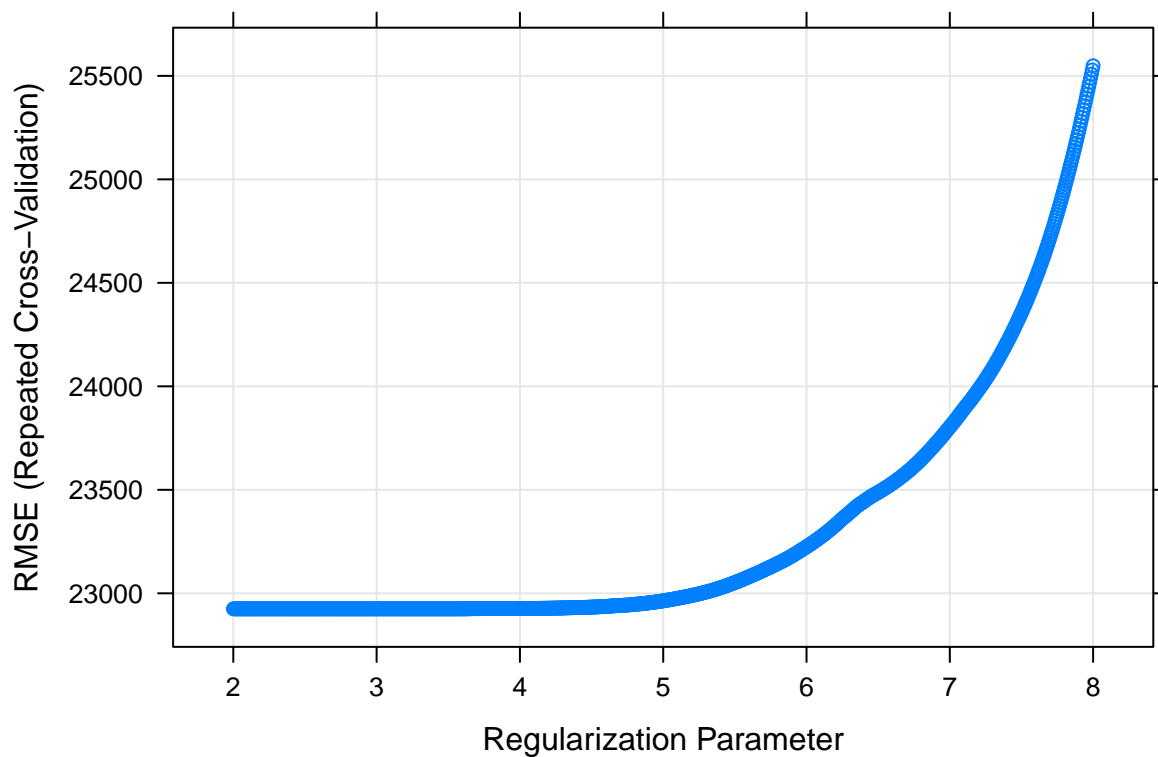
```
pred.lasso.1se = predict(cv.lasso, s = "lambda.1se", newx = xtest)
mse.lasso.1se = mean((ytest - pred.lasso.1se)^2)
mse.lasso.1se
```

```
## [1] 520300643
```

When fitting a Lasso model, the best tuning parameter for the minimum MSE rule is 75.0622912 and the test error is 4.7995097×10^8 . When the 1SE rule is applied, 29 predictors besides the intercept are included in the model.

LASSO by caret

```
set.seed(2023)
lasso.caret.min <- train(xtrain, ytrain,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(8, 2, length=1000))),
  trControl = ctrl1)
plot(lasso.caret.min, xTrans = log)
```



```
lasso.caret.min$bestTune
```

```
##      alpha  lambda
## 256      1 34.17627
```

```
coef(lasso.caret.min$finalModel, lasso.caret.min$bestTune$lambda)
```

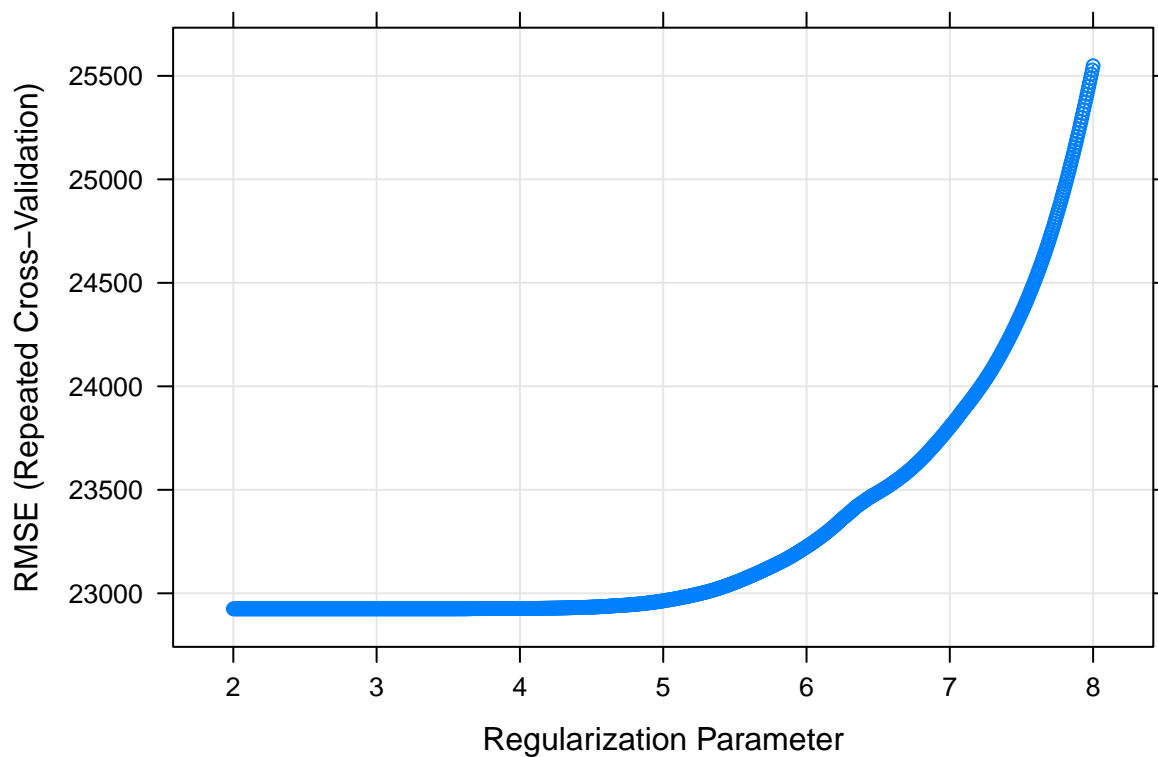
```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                                     s1
## (Intercept)                    -4.891889e+06
## gr_liv_area                      6.576019e+01
## first_flr_sf                     7.854992e-01
## second_flr_sf                    .
## total_bsmt_sf                   3.533373e+01
## low_qual_fin_sf                 -4.132375e+01
## wood_deck_sf                    1.179064e+01
## open_porch_sf                   1.574960e+01
## bsmt_unf_sf                     -2.089000e+01
## mas_vnr_area                    1.072139e+01
## garage_cars                     4.139513e+03
## garage_area                     8.020923e+00
## year_built                      3.241105e+02
## tot_rms_abv_grd                 -3.705645e+03
## full_bath                       -4.036604e+03
## overall_qualAverage             -4.924548e+03
## overall_qualBelow_Average      -1.260070e+04
## overall_qualExcellent           7.446458e+04
## overall_qualFair                -1.091887e+04
## overall_qualGood                1.218755e+04
## overall_qualVery_Excellent     1.337787e+05
## overall_qualVery_Good          3.793984e+04
```

```
## kitchen_qualFair      -2.556466e+04
## kitchen_qualGood      -1.785058e+04
## kitchen_qualTypical    -2.590666e+04
## fireplaces             1.087859e+04
## fireplace_quFair       -7.738271e+03
## fireplace_quGood       .
## fireplace_quNo_Fireplace 1.978771e+03
## fireplace_quPoor       -5.709901e+03
## fireplace_quTypical    -7.015803e+03
## exter_qualFair         -3.511041e+04
## exter_qualGood         -1.674606e+04
## exter_qualTypical      -2.116559e+04
## lot_frontage           1.007469e+02
## lot_area               6.044410e-01
## longitude              -3.366154e+04
## latitude               5.661433e+04
## misc_val               8.658075e-01
## year_sold              -5.939770e+02

pred.lasso.caret.min = predict(lasso.caret.min, s = lasso.caret.min$bestTune, newx = xtest)
mse.lasso.caret.min = mean((ytest - pred.lasso.caret.min)^2)
mse.lasso.caret.min

## [1] 479444173

ctrl2 = trainControl(method = "repeatedcv", number = 10, repeats = 5, selectionFunction = "oneSE")
set.seed(2023)
lasso.caret.1se <- train(xtrain, ytrain,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(8, 2, length=1000))),
  trControl = ctrl2)
plot(lasso.caret.1se, xTrans = log)
```



```
lasso.caret.1se$bestTune
```

```
##      alpha  lambda
## 670      1 410.7637
```

```
coef(lasso.caret.1se$finalModel, lasso.caret.1se$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                                     s1
## (Intercept)                    -3.897085e+06
## gr_liv_area                      6.091033e+01
## first_flr_sf                     9.502732e-01
## second_flr_sf                    .
## total_bsmt_sf                   3.629288e+01
## low_qual_fin_sf                 -3.504699e+01
## wood_deck_sf                    9.968877e+00
## open_porch_sf                   1.195645e+01
## bsmt_unf_sf                     -2.058462e+01
## mas_vnr_area                    1.300620e+01
## garage_cars                     3.479595e+03
## garage_area                     9.765822e+00
## year_built                      3.149071e+02
## tot_rms_abv_grd                 -2.497092e+03
## full_bath                       -1.367140e+03
## overall_qualAverage             -3.986739e+03
## overall_qualBelow_Average      -1.081091e+04
## overall_qualExcellent           8.736349e+04
## overall_qualFair                -8.718001e+03
## overall_qualGood                1.109250e+04
## overall_qualVery_Excellent     1.562134e+05
## overall_qualVery_Good          3.729801e+04
```



```
## kitchen_qualFair          -1.391060e+04
## kitchen_qualGood          -7.362904e+03
## kitchen_qualTypical       -1.618805e+04
## fireplaces                 8.135624e+03
## fireplace_quFair          -3.689784e+03
## fireplace_quGood           2.273316e+03
## fireplace_quNo_Fireplace   .
## fireplace_quPoor          -1.359037e+03
## fireplace_quTypical       -4.024181e+03
## exter_qualFair            -1.691749e+04
## exter_qualGood             .
## exter_qualTypical         -4.791414e+03
## lot_frontage               8.634683e+01
## lot_area                   5.911872e-01
## longitude                  -2.223730e+04
## latitude                   3.731517e+04
## misc_val                   2.964033e-01
## year_sold                  -1.583733e+02

pred.lasso.caret.1se = predict(lasso.caret.1se, s = lasso.caret.1se$bestTune, newx = xtest)
mse.lasso.caret.1se = mean((ytest - pred.lasso.caret.1se)^2)
mse.lasso.caret.1se

## [1] 496114286
```

We can also fit Lasso model using the `caret` package. The best tuning parameter for the minimum MSE rule is $\lambda = 410.7636622$ and the test error is `rmse.lasso.caret.min`. If we want to use the 1se rule, we can define a new resampling method that specifies `selectionFunction = "oneSE"`. With the 1se rule, there are 36 predictors included in the model.

elastic net

```
set.seed(2023)
enet.fit <- train(xtrain, ytrain,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                          lambda = exp(seq(7, -1, length = 200))),
                  trControl = ctrl1)
enet.fit$bestTune$alpha


## [1] 0.05

enet.fit$bestTune$lambda

## [1] 531.8609

myCol = rainbow(25)
myPar = list(superpose.symbol = list(col = myCol),
             superpose.line = list(col = myCol))

plot(enet.fit, par.settings = myPar)
```



```
pred.enet = predict(enet.fit, s = enet.fit$bestTune, newx = xtest)
mse.enet = mean((ytest - pred.enet)^2)
mse.enet
```

```
## [1] 480063606
```

When fitting an elastic net model, the selected tuning parameter for the minimum MSE rule is $\alpha = 0.05$ and $\lambda = 531.8608577$. It is possible to apply the 1SE rule to select the tuning parameters by using the resampling method of `ctrl2` where `selectionFunction = "oneSE"` is specified.

partial least squares by pls

```
set.seed(2023)
pls.fit = pls(sale_price~.,
              data = train,
              scale = TRUE,
              validation = "CV")
summary(pls.fit)
```

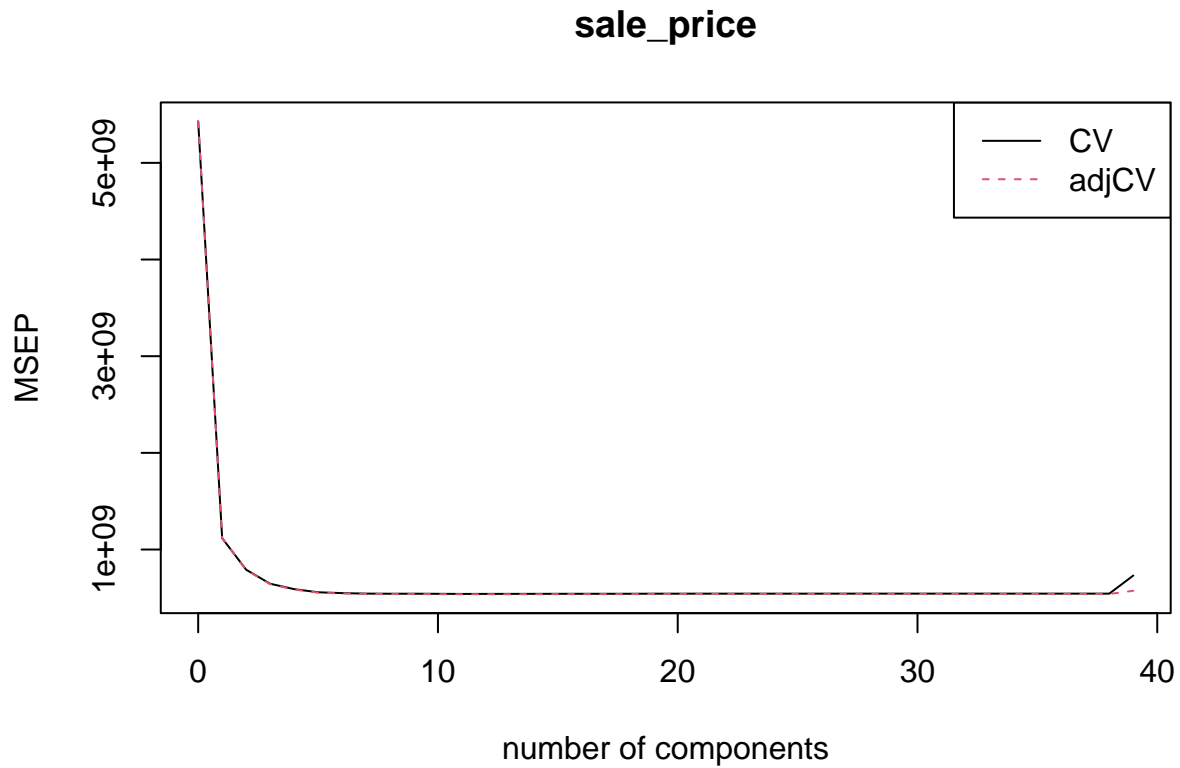
```
## Data:      X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           73685   33432   28131   25418   24296   23613   23430
## adjCV         73685   33427   28087   25329   24210   23534   23358
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
```

```

## CV      23324    23291    23300    23275    23252    23256    23261
## adjCV   23254    23222    23228    23203    23181    23183    23188
##      14 comps  15 comps  16 comps  17 comps  18 comps  19 comps  20 comps
## CV      23270    23275    23280    23283    23287    23304    23304
## adjCV   23196    23201    23206    23208    23213    23228    23228
##      21 comps  22 comps  23 comps  24 comps  25 comps  26 comps  27 comps
## CV      23309    23310    23310    23311    23312    23312    23315
## adjCV   23233    23234    23233    23234    23235    23235    23237
##      28 comps  29 comps  30 comps  31 comps  32 comps  33 comps  34 comps
## CV      23315    23315    23315    23316    23316    23316    23316
## adjCV   23238    23238    23238    23238    23238    23238    23238
##      35 comps  36 comps  37 comps  38 comps  39 comps
## CV      23316    23316    23316    23316    27032
## adjCV   23238    23238    23238    23238    23946
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          20.02    25.93    29.67    33.59    37.01    40.03    42.49
## sale_price  79.73    86.35    89.36    90.37    90.87    90.99    91.06
##      8 comps  9 comps  10 comps  11 comps  12 comps  13 comps  14 comps
## X          45.53    47.97    50.15    52.01    53.69    55.35    56.86
## sale_price  91.08    91.10    91.13    91.15    91.15    91.16    91.16
##      15 comps  16 comps  17 comps  18 comps  19 comps  20 comps
## X          58.64    60.01    62.18    63.87    65.26    67.10
## sale_price  91.16    91.16    91.16    91.16    91.16    91.16
##      21 comps  22 comps  23 comps  24 comps  25 comps  26 comps
## X          68.44    70.12    71.72    73.35    75.20    77.27
## sale_price  91.16    91.16    91.16    91.16    91.16    91.16
##      27 comps  28 comps  29 comps  30 comps  31 comps  32 comps
## X          78.97    80.10    81.83    83.55    84.39    86.34
## sale_price  91.16    91.16    91.16    91.16    91.16    91.16
##      33 comps  34 comps  35 comps  36 comps  37 comps  38 comps
## X          88.63    90.79    92.79    95.45    97.49    100.00
## sale_price  91.16    91.16    91.16    91.16    91.16    91.16
##      39 comps
## X          100.24
## sale_price  91.14

```

```
validationplot(pls.fit, val.type="MSEP", legendpos = "topright")
```



```
cv.mse = RMSEP(pls.fit)
ncomp.cv = which.min(cv.mse$val[1,,]) - 1
ncomp.cv
```

```
## 11 comps
##      11
```

```
pred.pls = predict(pls.fit, newdata = xtest, ncomp = ncomp.cv)
mse.pls = mean((ytest - pred.pls)^2)
mse.pls
```

```
## [1] 480106167
```

When fitting a partial least squares model using `pls`, the test error is 4.8010617×10^8 . There are 11 components included in the model.

partial least squares by caret

```
set.seed(2023)
pls.fit.caret = train(xtrain, ytrain,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:39),
  trControl = ctrl1,
  preProcess = c("center", "scale"))
pls.fit.caret$bestTune
```

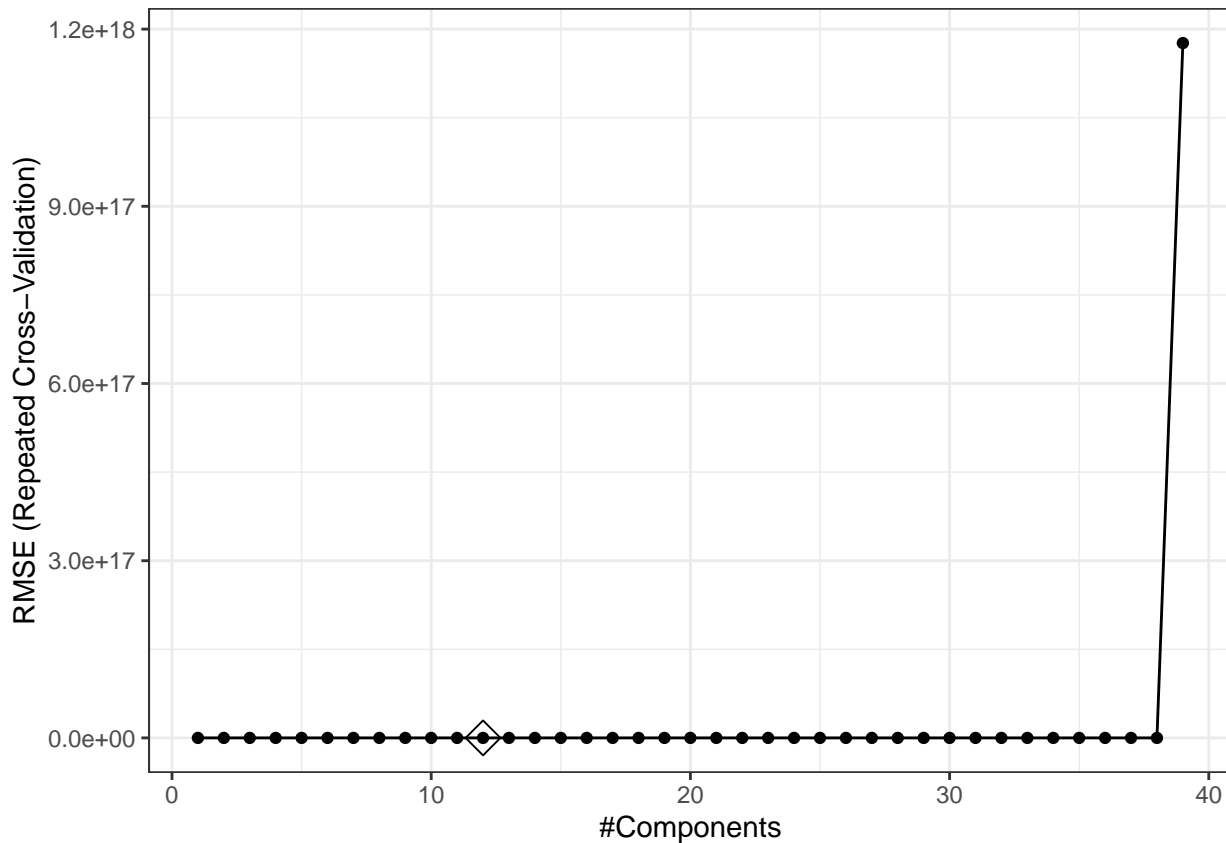
```
##      ncomp
## 12      12
```

```
pred.pls.caret = predict(pls.fit.caret, newdata = xtest)
mse.pls.caret = mean((ytest - pred.pls.caret)^2)
```

```
mse.pls.caret
```

```
## [1] 479655935
```

```
ggplot(pls.fit.caret, highlight = TRUE) + theme_bw()
```



We can also use `caret` package to fit a partial least squares model. We see that the number of components included in the model is different from what we got using `pls`.

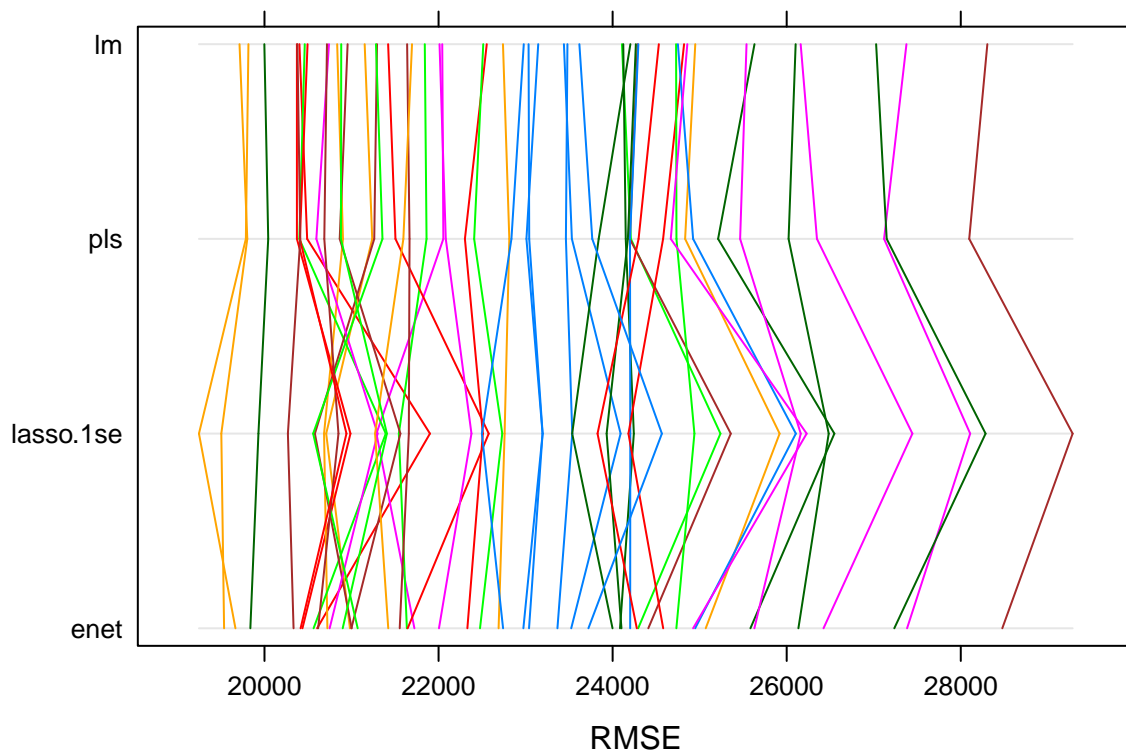
Comparing methods

```
resamp = resamples(list(lm = lm.fit, lasso.1se = lasso.caret.1se, enet = enet.fit, pls = pls.fit.caret))
summary(resamp)
```

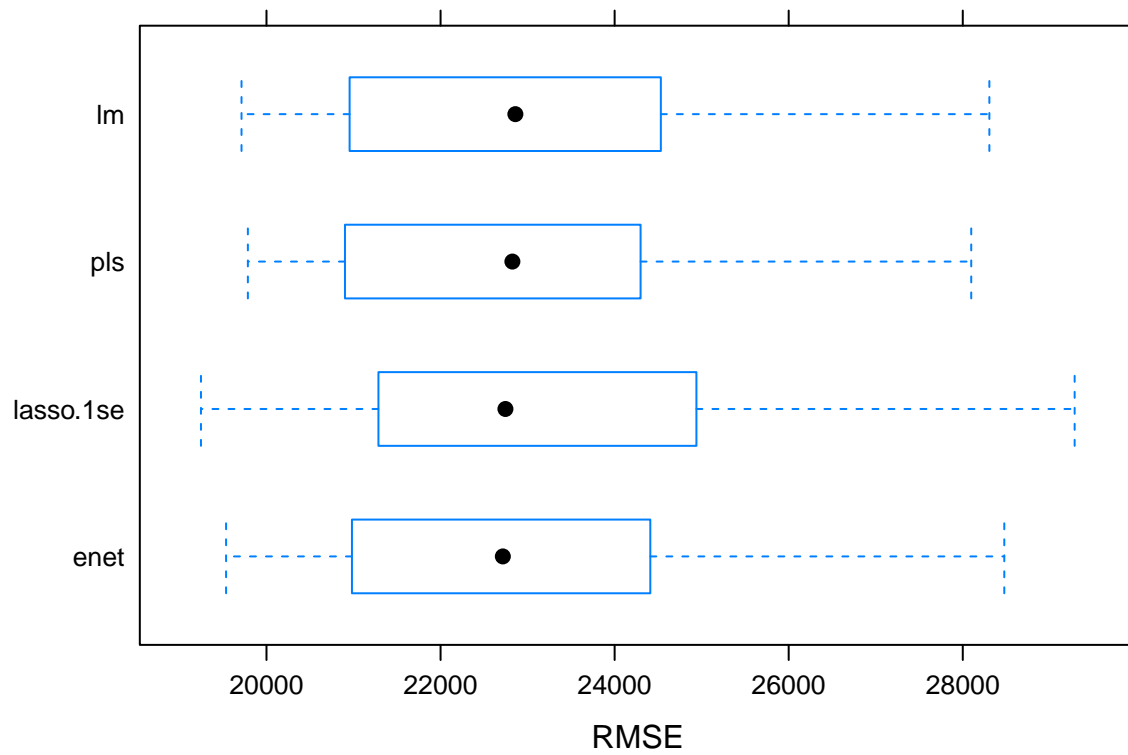
```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lm, lasso.1se, enet, pls
## Number of resamples: 50
##
## MAE
##           Min.   1st Qu.   Median     Mean  3rd Qu.     Max. NA's
## lm          13800.79 15933.82 16677.79 16706.93 17552.24 19577.64    0
## lasso.1se    13787.84 15756.68 16625.07 16650.15 17584.23 19299.33    0
## enet         13632.00 15849.10 16553.97 16627.25 17541.75 19493.36    0
## pls         13773.94 16021.17 16627.52 16703.52 17567.40 19558.61    0
##
```

```
## RMSE
##           Min.   1st Qu.   Median     Mean  3rd Qu.     Max. NA's
## lm       19713.17 21004.01 22859.59 22954.19 24472.24 28305.61    0
## lasso.1se 19247.12 21289.48 22745.25 23232.86 24846.25 29285.41    0
## enet      19535.62 20985.92 22715.59 22915.37 24380.82 28477.10    0
## pls      19785.94 20985.56 22825.16 22913.06 24276.12 28097.36    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean  3rd Qu.     Max. NA's
## lm       0.8659887 0.8911576 0.9014198 0.9035345 0.9181673 0.9431217    0
## lasso.1se 0.8530071 0.8922338 0.9019239 0.9018264 0.9146479 0.9399871    0
## enet      0.8636811 0.8928940 0.9018478 0.9039639 0.9176192 0.9430478    0
## pls      0.8645338 0.8922100 0.9025602 0.9037972 0.9180425 0.9433645    0
```

```
parallelplot(resamp, metric = "RMSE")
```



```
bwplot(resamp, metric = "RMSE")
```



By comparing across all models built, I would select elastic net for predicting the response because it has the smallest RMSE and MSE values. The adjusted R squares is also the second highest in the four models.